

A shallow introduction to Deep Learning with



DL4SCI 2020



**Evann Courdier**  
PhD, Deep  
Learning for  
Computer Vision

What to expect ?

# Webinar format

---

- 1 hour
- Prerequisites
  - Python
  - Gradient Based Machine Learning
  - A Numerical Computing Tool
- Slides & Notebooks

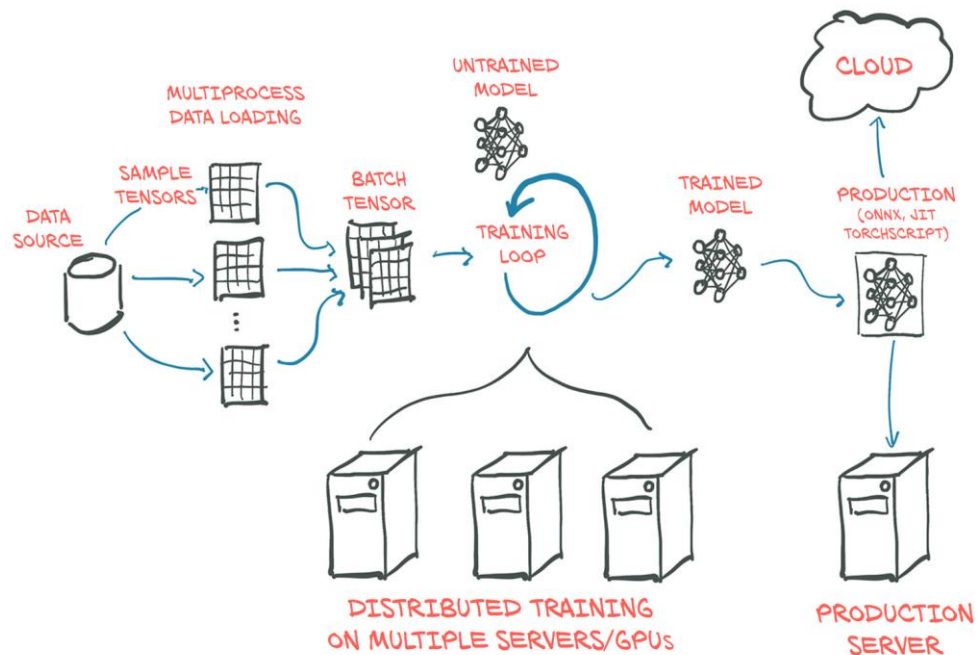
# PyTorch

---

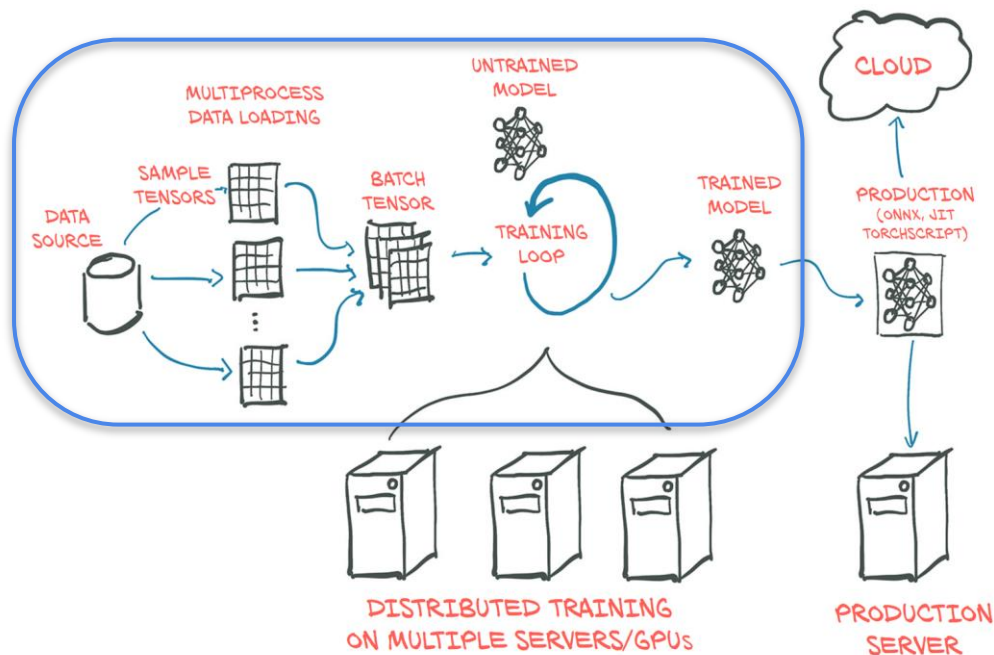
A library for scientific computing in Python, just like NumPy with:

- GPU support
- Automatic differentiation & Optimization algorithms
- All necessary tools for Deep Learning

# Deep Learning pipeline with PyTorch



# Deep Learning pipeline with PyTorch



# Program Skeleton

---

## INITIALISATION

Create model

Load data and prepare samples

Initialise optimisation and training HP

## TRAINING LOOP

Get a batch of samples and labels from the dataset

Move samples and labels to GPU

Compute the model's predictions

Compute the loss

Compute the gradients with backpropagation

Update the parameters



# It's all about Tensors

---

- Tensors are multi-dimensional arrays
- Very similar to NumPy for Tensor creation, indexing, masking

```
np.array([[1, 2, 3], [4, 5, 6]])
```

```
np.eye(2)
```

```
np.arange(1,5)
```

```
np.zeros(5)
```



```
torch.tensor([[1, 2, 3], [4, 5, 6]])
```

```
torch.eye(2)
```

```
torch.arange(1,5)
```

```
torch.zeros(5)
```



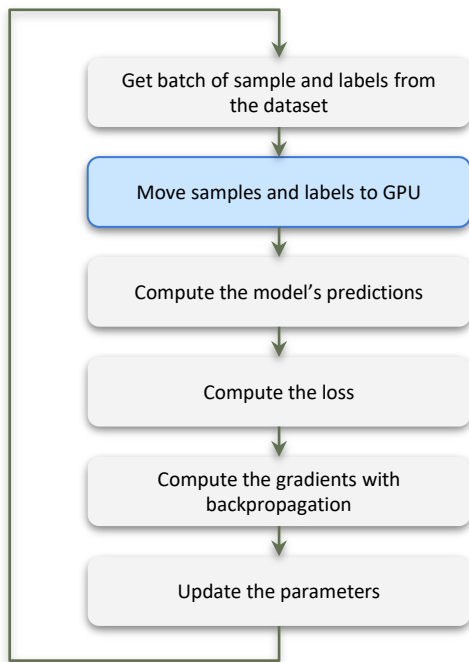
# Tensors - Recap

---

- Size and dimensions: `tensor.dim()` and `tensor.shape`
- Chain operations: `tensor.log().sum().exp()`
- In-place operations with underscore: `tensor.log_()`
- Reshape: `tensor.view(2,3)` and `tensor.view(-1, 3)`
- GPU  $\Leftrightarrow$  CPU: `torch.device` and `tensor.to(device)`

# Building the Training Loop

---



```
# TRAINING LOOP
```

```
# Loop through dataset to get batches of samples and labels
```

```
samples = samples.to(device)
```

```
labels = labels.to(device)
```



```
# compute predictions with model
```

```
# compute the loss
```



```
# compute gradients
```

```
# update model parameters
```


# Working with data: Dataset & Dataloaders



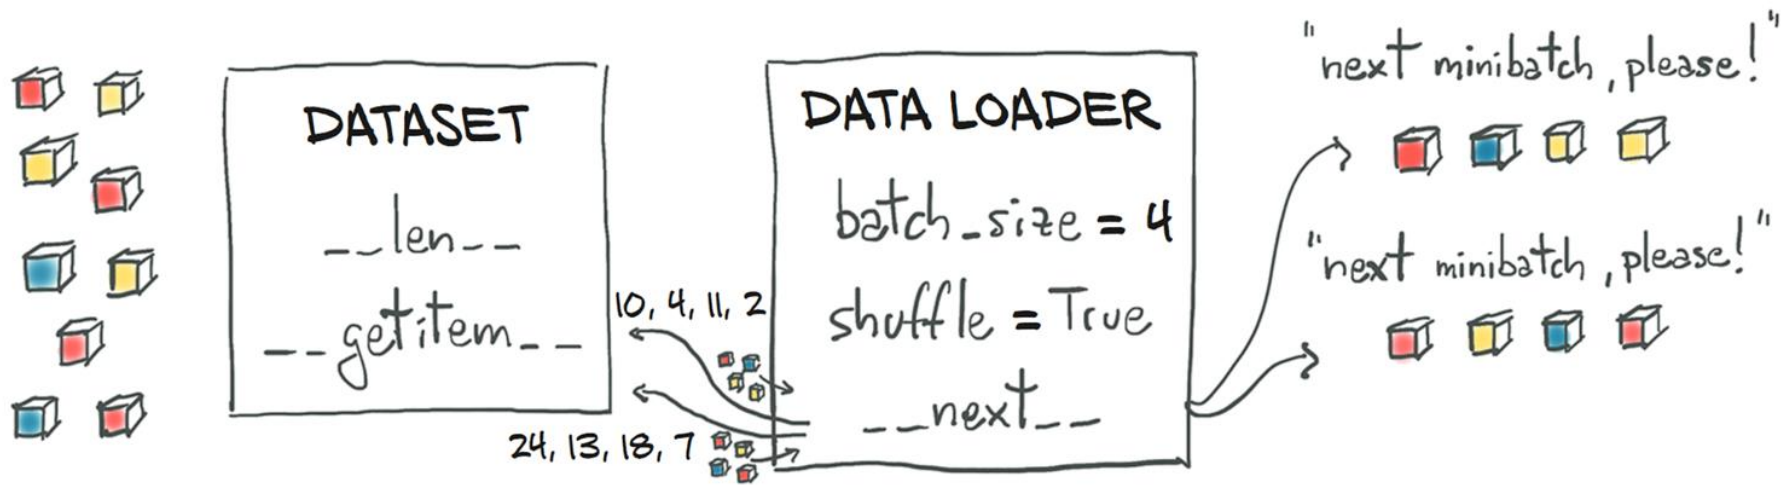
# Working with data: Dataset & Dataloaders



# Working with data: Dataset & Dataloaders



# Working with data: Dataset & Dataloaders

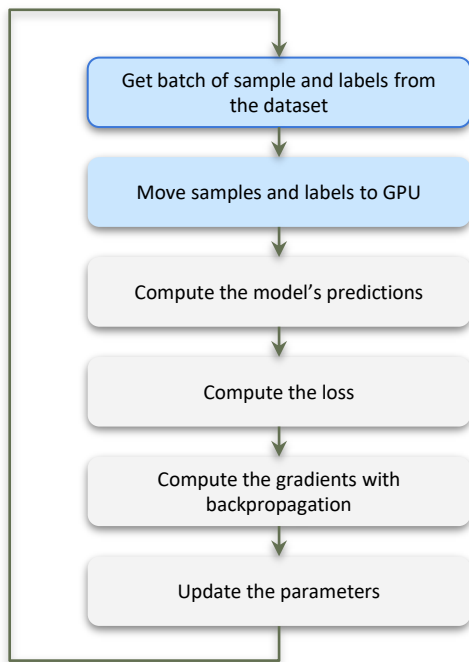






# Building the Training Loop

---



```
# TRAINING LOOP
```

```
for samples, labels in loader:  
    samples = samples.to(device)  
    labels = labels.to(device)  
    # compute predictions with model  
    # compute the loss  
    # compute gradients  
    # update model parameters
```

# Modules - Overview

---

- Help building reusable model components
- Manage model parameters
- PyTorch provides lots of built-in modules

# Modules - Managing Parameters

---

Modules help to:

- keep track of all parameters in your model.
- save/load your model
- reset all parameters gradients
- move all parameters to the gpu

# Modules - torch.nn

---

Whole library dedicated to Neural Network, including:

- Linear / Convolution / Recurrent Layers
- Activation Functions (ReLU, Tanh, ...)
- Loss Functions (MSE, CrossEntropy, ...)
- Pooling, Normalization, Dropout Layers

# Modules - torch.nn.Module

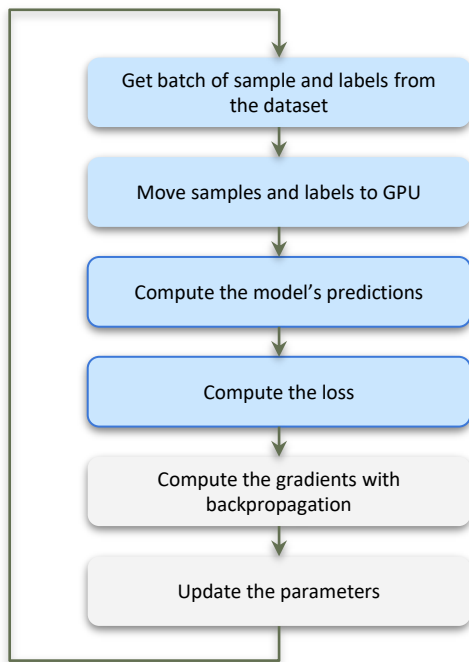
---

- A class you inherit from to create a Module
- It needs to implement two methods:
  - The `__init__` function: What are the components of your model
  - The `forward` function: How these components are connected



# Building the Training Loop

---



*# TRAINING LOOP*

```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    # compute gradients
    # update model parameters
```



# Computing gradients with Autograd

---

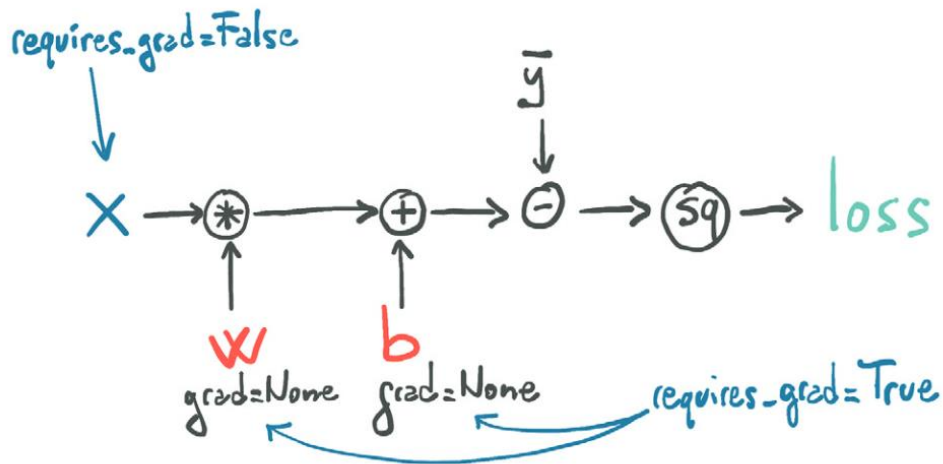
- Autograd: Automatic Differentiation package
- Each Tensor has a `requires_grad` boolean attribute
- Autograd creates a graph to record all operations during the computation
- Call `tensor.backward()` to compute all gradients automatically
- Gradients are accumulated into the `tensor.grad` attribute

# Computing gradients with Autograd

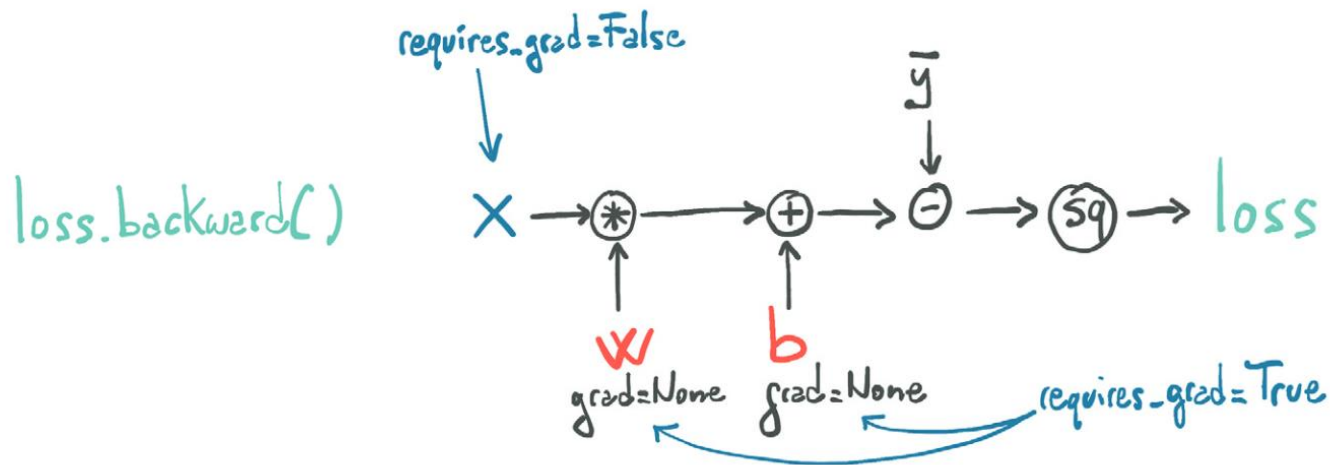
---

```
loss = (x * W + b - y) ** 2
```

# Computing gradients with Autograd




# Computing gradients with Autograd



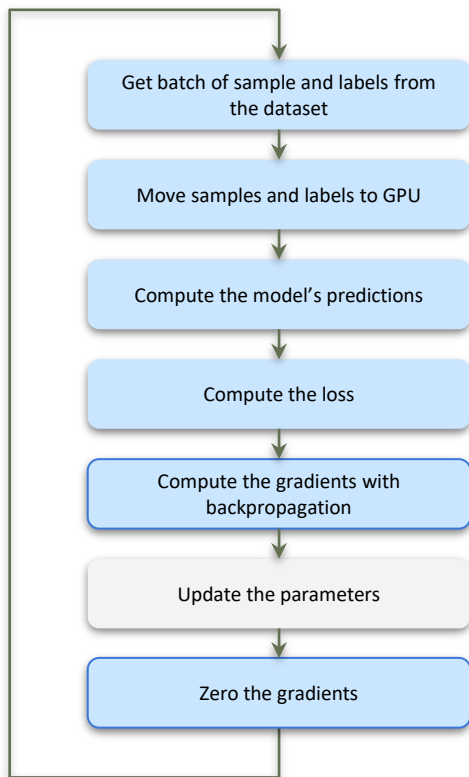
# Computing gradients with Autograd

`loss.backward()`





# Building the Training Loop



*# TRAINING LOOP*

```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    loss.backward()
    # update model parameters
    model.zero_grad()
```

# Optimizing parameters

---


- Use `torch.optim` submodule containing different optimizers
- The optimizer constructor takes a list of parameters
- A call to `optimizer.step()` updates the parameters

! Instead of `model.zero_grad()`,  
you can use `optimizer.zero_grad()`




# Optimizing parameters

---




# Optimizing parameters

---




# Optimizing parameters

---



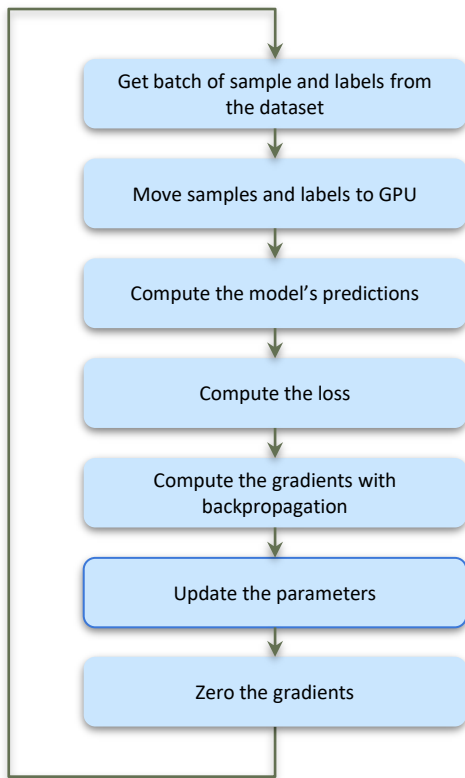
# Optimizing parameters

---





# Building the Training Loop




*# TRAINING LOOP*

```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    loss.backward()
    optimizer.step()
    model.zero_grad()
```

# Walkthrough: Building and Training LeNet on MNIST

---

- Building LeNet5 with `torch.nn.Module`
- Loading MNIST dataset with TorchVision
- Training for multiple epochs with a custom train function







# What you've still to discover

---

- Torchvision, Torchtext, Torchaudio
- Multi-GPU Distributed Training
- Quantisation & Pruning
- High-Performance with TorchScript + JIT
- Going to Production with ONNX, TorchElastic and TorchServe
- Organising your code with PyTorch Lightning



# Still wondering why PyTorch?

---



**Andrej Karpathy** ✓

@karpathy

Suivre



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

# Check our Material

---

Check [github.com/theevann/dl4sci-pytorch-webinar](https://github.com/theevann/dl4sci-pytorch-webinar) for:

- Notebooks presented during this webinar (live-notebooks)
- Detailed notebooks for offline study (offline-notebooks)

# Other References & Material

---

- Our notebooks at [github.com/theevann/dl4sci-pytorch-webinar/](https://github.com/theevann/dl4sci-pytorch-webinar/)
- Deep-Learning with PyTorch [e-book](#)
- Official tutorials at [pytorch.org/tutorials/](https://pytorch.org/tutorials/)
- A good tutorial on towardsdatascience at [bit.ly/38VgfaT/](https://bit.ly/38VgfaT/)
- Advanced EPFL Deep Learning Course at [fleuret.org/ee559/](https://fleuret.org/ee559/)

A shallow introduction to Deep Learning with



DL4SCI 2020