

A shallow introduction to Deep Learning with



DL4SCI 2020



Evann Courdier
PhD, Deep
Learning for
Computer Vision

What to expect ?

Webinar format

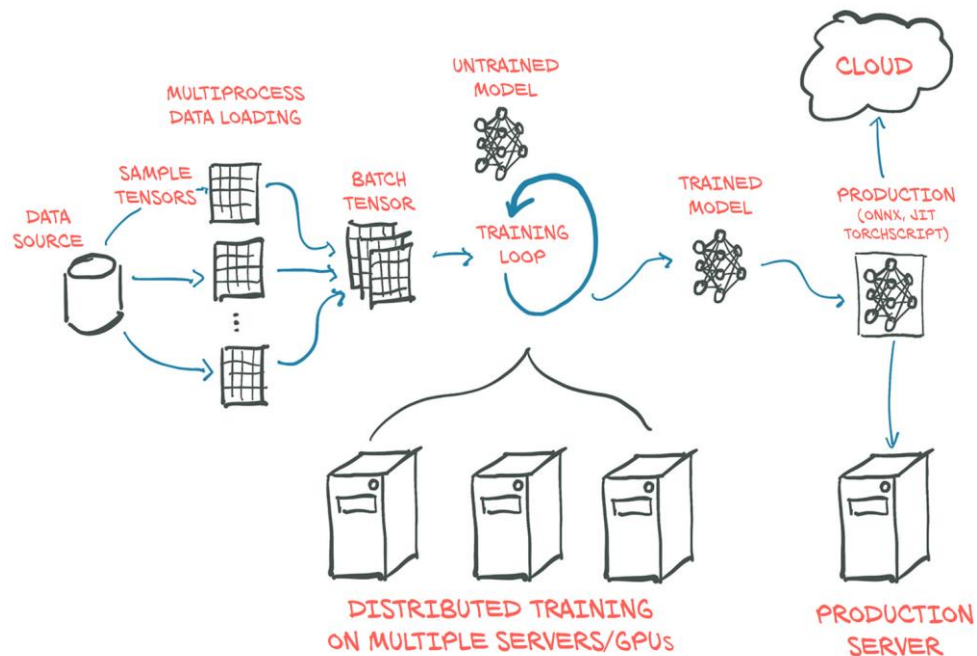
- 1 hour
- Prerequisites
 - Python
 - Gradient Based Machine Learning
 - A Numerical Computing Tool
- Slides & Notebooks

PyTorch

A library for scientific computing in Python, just like NumPy with:

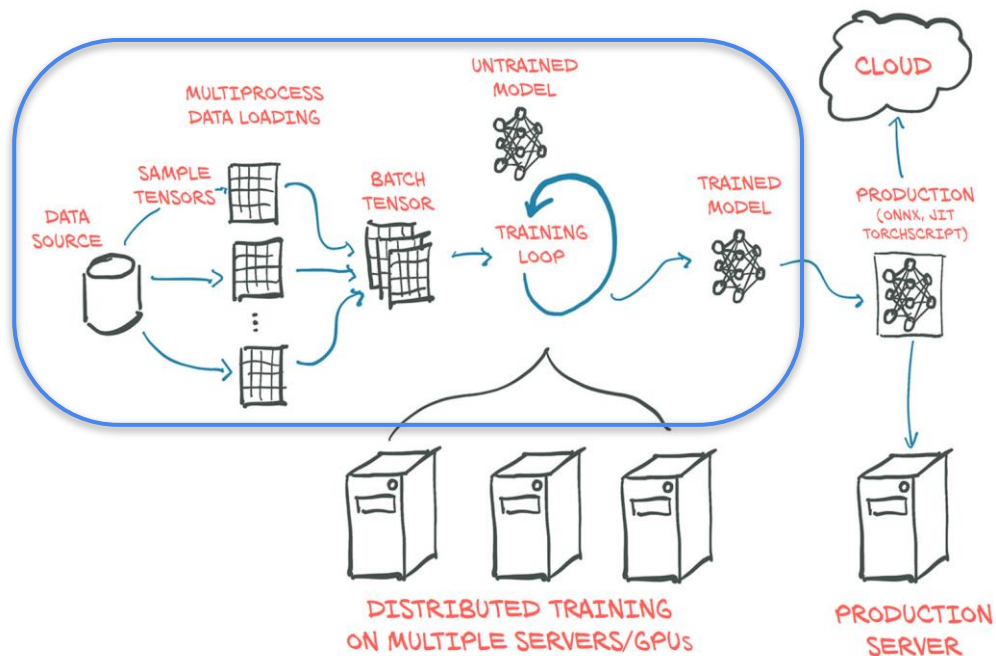
- GPU support
- Automatic differentiation & Optimization algorithms
- All necessary tools for Deep Learning

Deep Learning pipeline with PyTorch



All illustrations are taken from the Deep Learning with Pytorch e-book

Deep Learning pipeline with PyTorch



Program Skeleton

INITIALISATION

Create model

Load data and prepare samples

Initialise optimisation and training HP

TRAINING LOOP

Get a batch of samples and labels from the dataset

Move samples and labels to GPU

Compute the model's predictions

Compute the loss

Compute the gradients with backpropagation

Update the parameters

It's all about Tensors

- Tensors are multi-dimensional arrays
- Very similar to NumPy for Tensor creation, indexing, masking

```
np.array([[1, 2, 3], [4, 5, 6]])
```

```
np.eye(2)
```

```
np.arange(1,5)
```

```
np.zeros(5)
```



```
torch.tensor([[1, 2, 3], [4, 5, 6]])
```

```
torch.eye(2)
```

```
torch.arange(1,5)
```

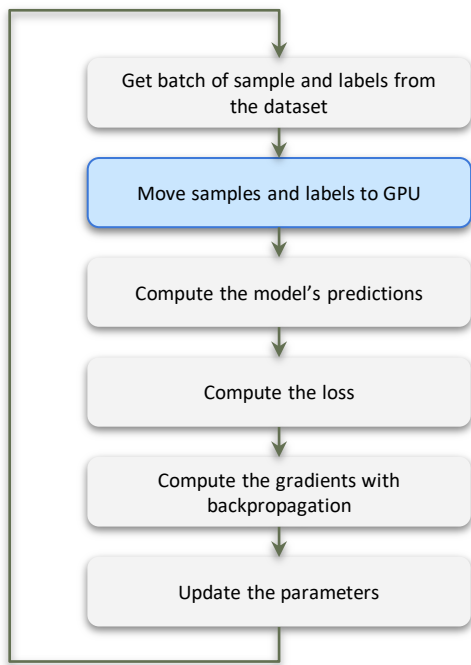
```
torch.zeros(5)
```



Tensors - Recap

- Size and dimensions: `tensor.dim()` and `tensor.shape`
- Chain operations: `tensor.log().sum().exp()`
- In-place operations with underscore: `tensor.log_()`
- Reshape: `tensor.view(2,3)` and `tensor.view(-1, 3)`
- GPU \Leftrightarrow CPU: `torch.device` and `tensor.to(device)`

Building the Training Loop



```
# TRAINING LOOP
```

```
# Loop through dataset to get batches of samples and labels
```

```
samples = samples.to(device)
```

```
labels = labels.to(device)
```

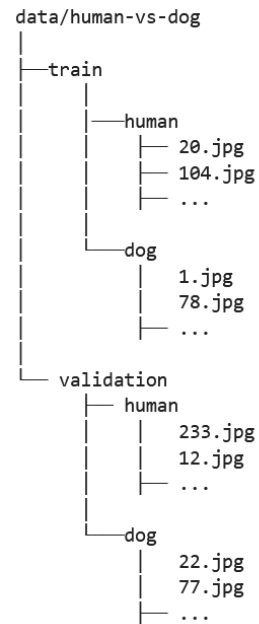

```
# compute predictions with model
```

```
# compute the loss
```



```
# compute gradients
```

```
# update model parameters
```


Working with data: Dataset & Dataloaders



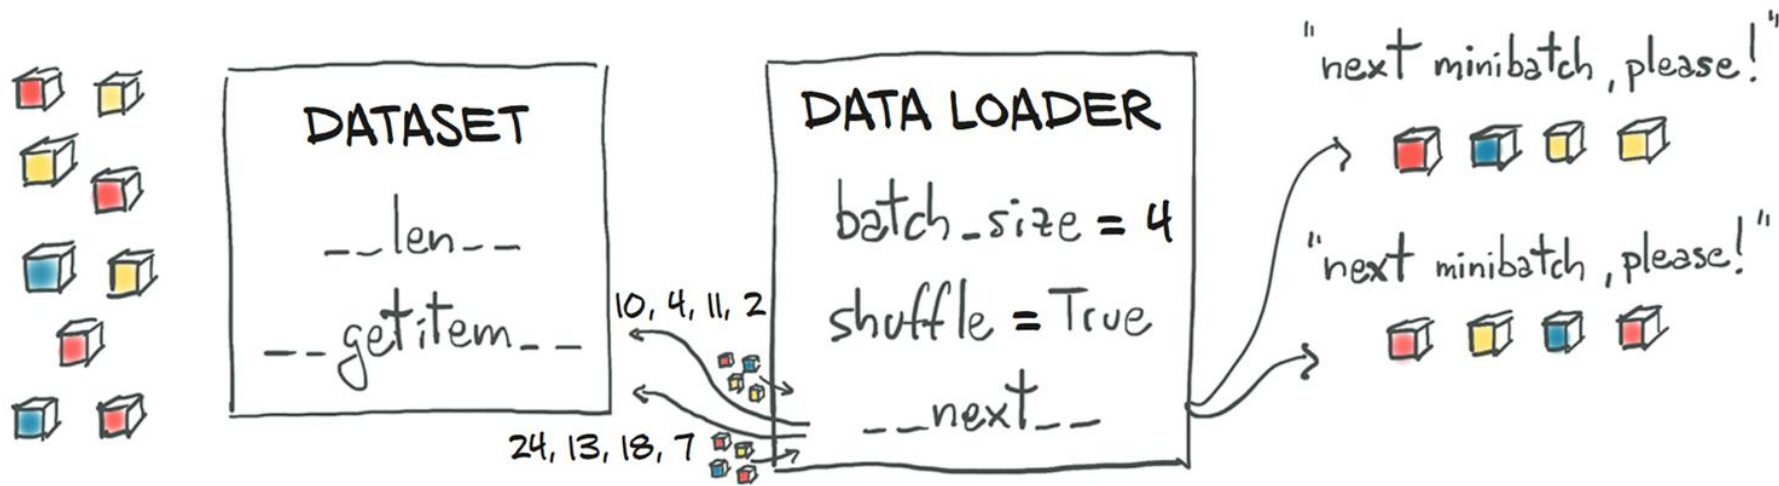
Working with data: Dataset & Dataloaders



Working with data: Dataset & Dataloaders

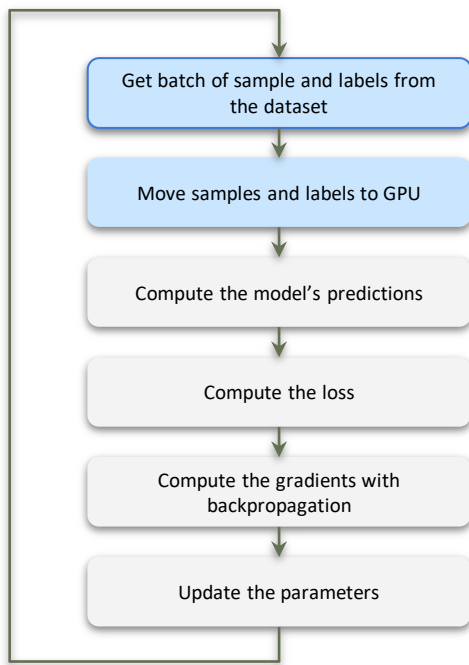


Working with data: Dataset & Dataloaders





Building the Training Loop



```
# TRAINING LOOP
```

```
for samples, labels in loader:  
    samples = samples.to(device)  
    labels = labels.to(device)  
    # compute predictions with model  
    # compute the loss  
    # compute gradients  
    # update model parameters
```

Modules - Overview

- Help building reusable model components
- Manage model parameters
- PyTorch provides lots of built-in modules

Modules - Managing Parameters

Modules help to:

- keep track of all parameters in your model.
- save/load your model
- reset all parameters gradients
- move all parameters to the gpu

Modules - torch.nn

Whole library dedicated to Neural Network, including:

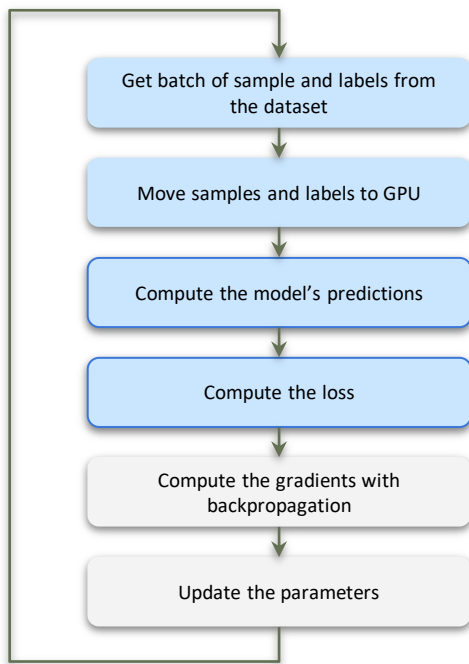
- Linear / Convolution / Recurrent Layers
- Activation Functions (ReLU, Tanh, ...)
- Loss Functions (MSE, CrossEntropy, ...)
- Pooling, Normalization, Dropout Layers

Modules - torch.nn.Module

- A class you inherit from to create a Module
- It needs to implement two methods:
 - The `__init__` function: What are the components of your model
 - The `forward` function: How these components are connected



Building the Training Loop



TRAINING LOOP

```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    # compute gradients
    # update model parameters
```

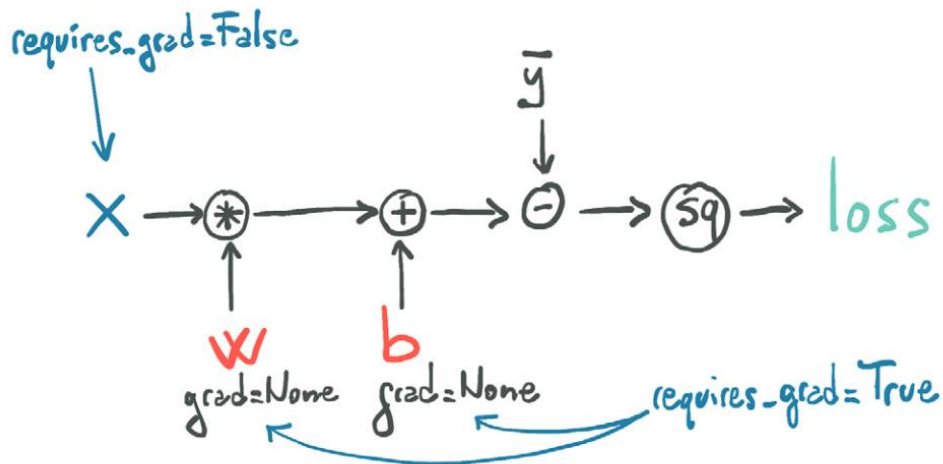

Computing gradients with Autograd

- Autograd: Automatic Differentiation package
- Each Tensor has a `requires_grad` boolean attribute
- Autograd creates a graph to record all operations during the computation
- Call `tensor.backward()` to compute all gradients automatically
- Gradients are accumulated into the `tensor.grad` attribute

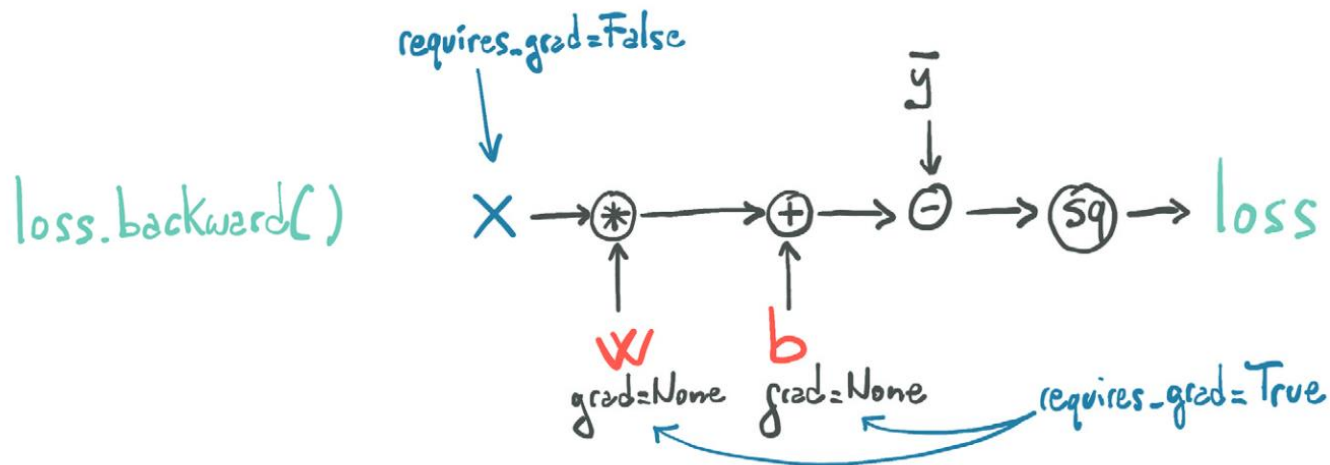
Computing gradients with Autograd

```
loss = (x * W + b - y) ** 2
```


Computing gradients with Autograd



Computing gradients with Autograd

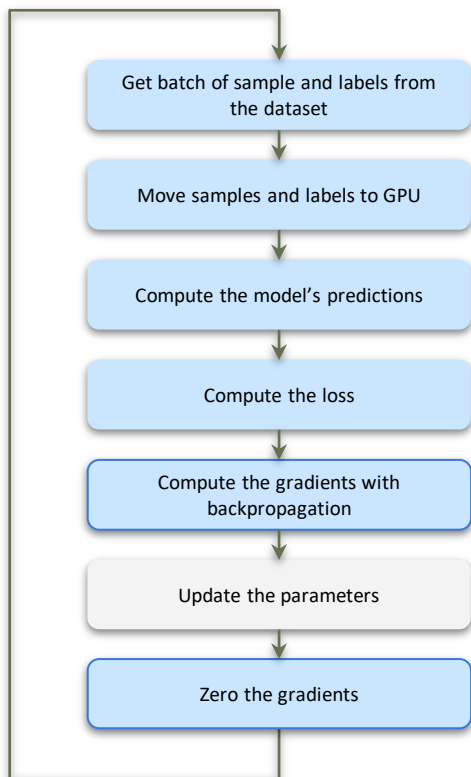


Computing gradients with Autograd





Building the Training Loop



TRAINING LOOP


```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    loss.backward()
    # update model parameters
    model.zero_grad()
```

Optimizing parameters


- Use `torch.optim` submodule containing different optimizers
- The optimizer constructor takes a list of parameters
- A call to `optimizer.step()` updates the parameters

! Instead of `model.zero_grad()`,
you can use `optimizer.zero_grad()`


Optimizing parameters




Optimizing parameters



Optimizing parameters

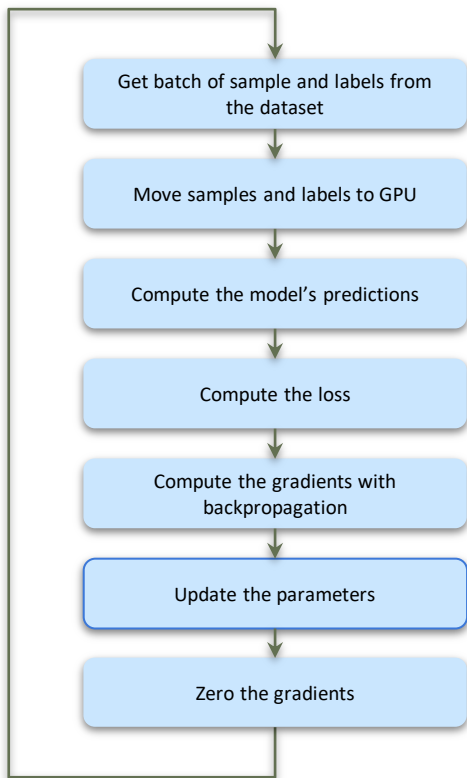


Optimizing parameters





Building the Training Loop




TRAINING LOOP

```
for samples, labels in loader:
    samples = samples.to(device)
    labels = labels.to(device)
    predictions = model(samples)
    loss = loss_fn(predictions, labels)
    loss.backward()
    optimizer.step()
    model.zero_grad()
```

Walkthrough: Building and Training LeNet on MNIST

- Building LeNet5 with `torch.nn.Module`
- Loading MNIST dataset with TorchVision
- Training for multiple epochs with a custom train function





What you've still to discover

- Torchvision, Torchtext, Torchaudio
- Multi-GPU Distributed Training
- Quantisation & Pruning
- High-Performance with TorchScript + JIT
- Going to Production with ONNX, TorchElastic and TorchServe
- Organising your code with PyTorch Lightning



Still wondering why PyTorch?



Andrej Karpathy ✓

@karpathy

Suivre



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

Check our Material

Check github.com/theevann/dl4sci-pytorch-webinar for:

- Notebooks presented during this webinar (live-notebooks)
- Detailed notebooks for offline study (offline-notebooks)

Other References & Material

- Our notebooks at github.com/theevann/dl4sci-pytorch-webinar/
- Deep-Learning with PyTorch [e-book](#)
- Official tutorials at pytorch.org/tutorials/
- A good tutorial on towardsdatascience at bit.ly/38VgfaT/
- Advanced EPFL Deep Learning Course at fleuret.org/ee559/

A shallow introduction to Deep Learning with



DL4SCI 2020